

# SQL Views

**SQL views** are virtual tables that are created using a **SELECT** statement in SQL. A view is a database object that acts as a filter to the data stored in one or more tables. It is a logical representation of data in a database that can be used to simplify the complexity of data and enhance security.

## Syntax

Create view v\_name

As

Select \* column 1 and column 2

From table\_name

Where condition;

**Advantages to using views:**

- views can simplify complex queries and make them easier to read and understand.
- Views can also be used to enhance security by restricting access to certain data.
- views can be used to hide sensitive data from certain users or applications.

## Types of Views in SQL Server:

- 1- *Simple View* is a view created from a single table and does not use complex SQL features.
  - Is based on only one table
  - **Does not use:**
  - JOIN
  - GROUP BY
  - HAVING
  - Aggregate functions (SUM, AVG, COUNT, etc.)

## Syntax

```
CREATE VIEW name_View AS
```

```
SELECT column1, column2
```

```
FROM table_name
```

WHERE condition;

## Real-life use cases

### 1 User Access & Security (Most Common Use)

#### Problem:

You don't want users to see sensitive columns (salary, password, SSN).

#### Solution (Simple View):

```
CREATE VIEW Employee_Public AS  
SELECT EmployeeID, Name, Department  
FROM Employee;
```

### 2 Active Records Only (Filtering Data)

#### Problem:

Applications frequently need **only active records**.

#### Solution:

```
CREATE VIEW ActiveCustomers AS  
SELECT CustomerID, CustomerName, Email  
FROM Customers  
WHERE Status = 'Active';
```

2- Indexed View is a view that has a clustered index created on it, which means:

the result of the view is stored physically on disk, just like a table.

Indexed Views are used to:

- Improve performance
- Speed up JOINs and aggregate queries
- Reduce repeated calculations
- Optimize reporting queries

Key Difference from Normal Views

Feature	Normal View	Indexed View
Stores data	 No (virtual)	 Yes (physical)

Performance	Normal	Faster
Uses index	No	Yes
SQL Server support		(only SQL Server)

## == When to Use Indexed Views?

- ✓ Large tables
- ✓ Heavy aggregation (SUM, COUNT, AVG)
- ✓ Frequently executed queries
- ✓ Reporting systems

## == When NOT to Use Them?

- Tables with frequent INSERT/UPDATE/DELETE
- Small datasets
- When schema changes often

## Real-life use cases

### 1 E-Commerce Website (Fast Sales Dashboards)

#### Problem:

Managers frequently check **total sales per product/category**.  
Running SUM() and GROUP BY on millions of rows is slow.

Indexed View Solution:

- ✓ Dashboard loads instantly
- ✓ Aggregates are already stored
- ✓ Very common in large e-commerce systems

### 2 Banking System (Account Balance Summaries)

**Problem:**

Balance is calculated from thousands of transactions every time.

**Indexed View Solution:**

- View stores **total deposits and withdrawals per account**
- Clustered index keeps data ready

- ✓ Faster balance checks
- ✓ Used in ATMs and online banking
- ✓ Avoids recalculating totals repeatedly

3--University / Education System (Results & GPA)

**Problem:**

Calculating GPA for thousands of students repeatedly is expensive.

**Indexed View Solution:**

- Pre-calculated **GPA per student**
- Indexed for fast access

- ✓ Faster student portals
- ✓ Efficient result processing
- ✓ Stable reporting

3- Partitioned View a view that combines data from multiple tables that have the same structure, usually to make them behave like one large table.

Partitioned Views are used to:

- Manage very large tables
- Improve performance
- Separate data logically (by year, region, department, etc.)
- Support scalability

## How does it work?

- Data is split across multiple tables
- Each table contains different rows
- A CHECK constraint defines which rows belong to each table
- A view combines them using UNION ALL

### Example Scenario:

```
CREATE TABLE Orders_2023 (
    OrderID INT PRIMARY KEY,
    OrderDate DATE,
    Amount DECIMAL(10,2),
    CHECK (YEAR(OrderDate) = 2023)
);

CREATE TABLE Orders_2024 (
    OrderID INT PRIMARY KEY,
    OrderDate DATE,
    Amount DECIMAL(10,2),
    CHECK (YEAR(OrderDate) = 2024)
);
```

### Create the Partitioned View

```
CREATE VIEW Orders_All
AS
SELECT * FROM Orders_2023
UNION ALL
SELECT * FROM Orders_2024;
```

### Partitioned View vs Indexed View

Feature	Partitioned View	Indexed View
Purpose	Split large data	Speed up queries
Data storage	Separate tables	Stored via index
Uses UNION ALL	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No
Performance gain	Scalability	Query speed

## Real-life use cases

### 1 Banking System (Transactions by Year)

#### Problem:

A bank has **millions of transactions**. One table becomes very slow.

#### Solution using Partitioned View:

- Transactions\_2023
- Transactions\_2024
- Transactions\_2025

Each table stores one year of data.

```
CREATE VIEW AllTransactions AS
SELECT * FROM Transactions_2023
UNION ALL
SELECT * FROM Transactions_2024
UNION ALL
SELECT * FROM Transactions_2025;
```

### 2 University / College Database (Students by Batch)

#### Problem:

Student records grow every year.

#### Solution:

- Students\_2022
- Students\_2023
- Students\_2024

```
CREATE VIEW AllStudents AS
SELECT * FROM Students_2022
UNION ALL
SELECT * FROM Students_2023
UNION ALL
SELECT * FROM Students_2024;
```

### E-Commerce System (Orders by Region)

#### **Problem:**

Orders are stored globally → slow queries.

#### **Solution:**

- Orders\_US
- Orders\_EU
- Orders\_ASIA

```
CREATE VIEW AllOrders AS
SELECT * FROM Orders_US
UNION ALL
SELECT * FROM Orders_EU
UNION ALL
SELECT * FROM Orders_ASIA;
```

. Can We Use DML (INSERT, UPDATE, DELETE) on Views? Yes we can

Which types of views allow DML operations?

=simple type

= View with check option

= Join View

#### **What are the restrictions or limitations when performing DML on a view?**

##### **1. View must be updatable**

- a. Only **simple views** (based on one table, no aggregates, no DISTINCT, no GROUP BY, no UNION) are fully updatable.
- b. Views that are **complex** (joins, aggregates, distinct, union, computed columns) are either partially updatable or non-updatable.

##### **2. Columns must belong to a single base table for updates**

- a. If the view is a join, you can update **only the columns from one underlying table** at a time.

### 3. Cannot modify derived/computed columns

```
CREATE VIEW vw_Salary
AS
SELECT EmployeeID, Salary * 1.1 AS NewSalary
FROM Employees;
```

- a. ✗ Cannot insert or update NewSalary because it's computed.

### 4. Cannot perform DML that violates view constraints

- a. If the view has a WITH CHECK OPTION, you cannot insert/update rows that would be excluded from the view.

**Give at least one real-life example where updating a view is useful (e.g., HR system, e-commerce orders, etc.)**

#### Scenario: HR System – Updating Employee Contact Info

##### Problem

- HR has a large **Employees** table with sensitive columns (Salary, Social Security Number, etc.).
- HR staff need to **update only contact information** (phone, email) without seeing or touching sensitive data.

##### Step 1: Create a Simple View for Safe Updates

```
CREATE VIEW vw_EmployeeContacts
AS
SELECT EmployeeID, Name, Phone, Email
FROM Employees;
```

##### Step 2: Update via the View

```
UPDATE vw_EmployeeContacts
SET Phone = '987-654-3210',
    Email = 'new.email@example.com'
WHERE EmployeeID = 101;
```

## **2. How Can Views Simplify Complex Queries?**

- Explain how a View can help simplify JOIN-heavy queries.

## **1 Problem: JOIN-heavy Queries Are Complex**

When a query involves **many tables**, joins, and conditions, it can become long, hard to read, and error-prone.

- Multiple joins make the query **long and difficult to maintain**.
- Rewriting it for reports or analytics can be tedious.

## **2 Solution: Use a View to Encapsulate Joins**

- A view can **predefine the joins**, so users don't have to rewrite them every time.
- Makes queries **shorter, readable, and maintainable**.

**Create an example view that joins at least two of your banking tables, such as: o Customer + Account o Account + Transaction**

### **1 View: Customer + Account**

```
CREATE VIEW vw_CustomerAccounts
AS
SELECT
    c.CustomerID, c.FullName, c.Email, a.AccountID, a.AccountType, a.Balance, a.OpenDate
FROM Customers c
JOIN Accounts a
ON c.CustomerID = a.CustomerID;
```

### **2 View: Account + Transaction**

```
CREATE VIEW vw_AccountTransactions
AS
SELECT
    a.AccountID, a.AccountType, a.Balance AS CurrentBalance, t.TransactionID,
```

t.TransactionDate, t.Amount, t.TransactionType

FROM Accounts a

JOIN Transactions t

ON a.AccountID = t.AccountID;